MapShift: Controlled Post-Intervention Evaluation for Embodied Planning

Anonymous Author(s)
Affiliation
Address
email

Abstract

Embodied agents are often evaluated in the same environment in which they were explored or trained, making it difficult to assess whether a learned world model supports planning after the world changes. Existing evaluations can conflate memory of the explored environment, belief update after change, and planning under the updated belief. We introduce MapShift, an executable benchmark for controlled post-intervention evaluation (CEP): an agent explores a base environment without task reward, the environment is modified by a controlled metric, topology, dynamics, or semantic intervention, and the agent is evaluated on post-intervention planning, inference, and adaptation tasks. The contribution is measurement infrastructure: matched base/intervened pairs, family-wise estimands, severity ladders, invariant validators, benchmark-health gates, protocol-comparison tooling, and reproducible artifact generation. In a reference study, health gates pass with zero fatal leakage, zero task rejections, perfect reference solvability, no intervention-validator failures, and no severity-magnitude failures. The main empirical finding is that protocol choices change conclusions about world-model competence: no-exploration changes all four family-wise method orderings, and a deterministic mechanism diagnostic shows that same-environment evaluation can rank stale memory too highly while CEP separates stale reuse, local heuristics, and explicit belief update.

1 Introduction

World models are useful only insofar as their learned structure remains actionable when the world does not stay fixed. In embodied evaluation, however, agents are commonly trained, explored, and evaluated under tightly coupled environment assumptions. Even when environments are procedurally generated or randomized, the evaluation often asks whether a system performs well under a sampled distribution, not whether the system can use reward-free exploration to reason about a specific environment after a controlled change. This distinction matters because protocol design determines which competence is being measured: stale reuse of an explored map, explicit belief update after a change, or planning under the updated belief. Embodied planning combines partial observability and sequential decision-making with structured spatial state [Kaelbling et al., 1998, Sutton and Barto, 2018]; failures may arise from brittle geometry, missing connectivity, incorrect transition dynamics, or shifted cue meanings.

We study this question as an evaluation problem rather than as a new modeling proposal. The central object is controlled post-intervention evaluation (CEP): sample a base environment e, let an agent explore e without task reward for a fixed budget T_{exp}, apply an intervention $I_f(\sigma)$ to obtain a changed environment e', and evaluate the agent on tasks sampled from $Q(e, e')$. The family f identifies what changed, and the severity σ controls how strongly it changed. We use the term "counterfactual" only

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

in this operational matched-pair sense: the generator seed and motif are held fixed while a declared environment factor is changed. The protocol is therefore closer to controlled out-of-distribution evaluation than to unit-level causal identification from observational data. Under this protocol, a method can fail despite strong same-environment performance if its internal representation is not useful after the intervention. Figure 1 summarizes the evaluation flow.

MapShift is a controlled benchmark implementation of this protocol. It is designed as executable measurement infrastructure: it provides matched intervention pairs, family-wise estimands, explicit severity ladders, motif-level splits, reproducible task generation, benchmark health reports, protocol comparisons, and paper-facing artifact builders. MapShift is intended for researchers comparing post-intervention robustness of embodied planners, evaluating new world-model architectures against intervention-family stress profiles, or auditing whether evaluation conclusions depend on protocol assumptions. The substrate is deliberately small so that intervention isolation, task solvability, leakage, and weak-baseline saturation can be verified directly. The release should therefore be read as a calibration suite and reference standard, not as a scale or realism benchmark: it specifies what an intervention wrapper should prove, log, and report before the same protocol is ported to richer but harder-to-audit environments such as MiniGrid, ProcTHOR, or Habitat.

The paper is organized around one central evaluation-science question: when do protocol choices change conclusions about world-model competence? We answer it through family-wise reporting, no-exploration and same-environment controls, and a deterministic mechanism diagnostic that separates stale maps, local heuristics, and explicit belief update. More broadly, MapShift targets a recurring evaluation failure in embodied AI: aggregate task success can obscure whether an agent updated its world model, reused stale state, or merely exploited local heuristics.

Our contributions are:

1. We formalize a controlled post-intervention evaluation protocol for measuring embodied competence after reward-free exploration.
2. We release MapShift, a benchmark with eight motifs, four intervention families, four severity levels, three task classes, and $T_{\text{exp}} = 800$.
3. We make the validation contract for post-intervention wrappers executable: matched-pair construction, invariant checks, leakage/coverage/rejection/solvability/saturation health gates, raw records, and protocol-comparison artifacts are generated with the benchmark.
4. We define a reproducible study artifact that compares privileged planning reference, same-environment, diagnostic, heuristic, classical belief-update, pretrained graph, memory, relational, and structured-dynamics baselines across model seeds and protocol variants.

2 Related work

Reinforcement-learning and embodied benchmarks. Benchmark environments such as ALE, Gym, DeepMind Control/Lab, Procgen, and MiniGrid support reproducible agent research [Bellemare et al., 2013, Brockman et al., 2016, Tassa et al., 2018, Beattie et al., 2016, Cobbe et al., 2020, Chevalier-Boisvert et al., 2023]; Crafter and MineDojo add open-ended skills [Hafner, 2022, Fan et al., 2022]; and Habitat, AI2-THOR, ProcTHOR, VirtualHome, and ALFRED add richer embodied structure [Savva et al., 2019, Kolve et al., 2017, Deitke et al., 2022, Puig et al., 2018, Shridhar et al., 2020]. MapShift is complementary: it is not a realism claim, but a reference implementation for testing whether reward-free environment knowledge remains useful after controlled intervention. A wrapper for MiniGrid, ProcTHOR, or Habitat could instantiate CEP, but would still need to prove matched-pair invariants, rejection behavior, solvability, weak-baseline saturation, raw-record reproducibility, and protocol comparisons. MapShift makes that measurement contract cheap to audit before porting it to richer simulators.

World models and planning under partial observability. World-model and model-based RL methods learn predictive structure for planning and control [Ha and Schmidhuber, 2018, Hafner et al., 2020, Schrittwieser et al., 2020], while reward-free exploration separates information gathering from downstream reward optimization [Jin et al., 2020, Sekar et al., 2020, Mendonca et al., 2021]. MapShift asks a narrower diagnostic question under partial observability [Kaelbling et al.,

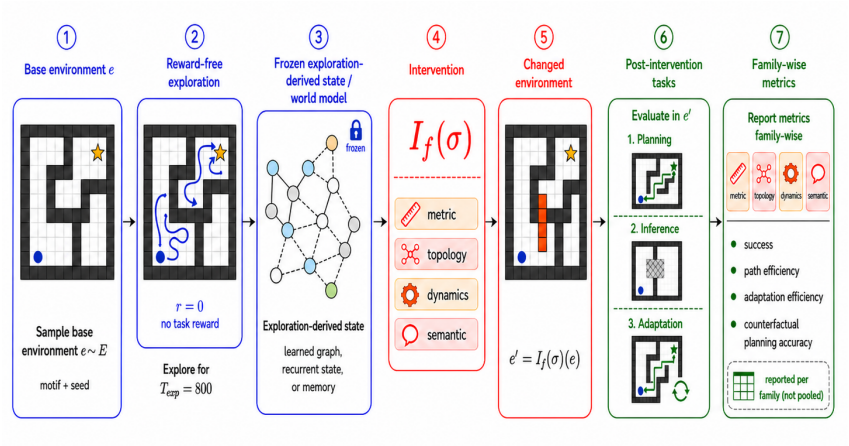

Figure 1: Controlled post-intervention evaluation protocol. An agent first explores a base environment without task reward, producing a fixed exploration-derived memory or world model. A controlled intervention then changes the environment along one family and severity level. The agent is evaluated only after the intervention on planning, inference, and adaptation tasks, and results are reported by intervention family to expose intervention-specific failure modes.

1998]: whether the learned representation still supports update, inference, and planning after geometry, topology, dynamics, or semantics change.

Reward-free world-model evaluation. WorldTest/AutumnBench also separates reward-free interaction from a scored derived test phase [Warrier et al., 2025]. MapShift differs by making the measurement contract executable: declared intervention families, matched pairs, invariant validators, explicit severity ladders, family-wise estimands, and health gates for leakage, rejection, solvability, weak-baseline saturation, and severity response.

Causal and counterfactual reinforcement learning. Causal reinforcement-learning work studies confounding, interventions, and counterfactual reasoning in sequential decision processes [Lu et al., 2018, Buesing et al., 2019, Rezende et al., 2020]. MapShift does not attempt to identify causal effects from observational trajectories, nor does it assume that the agent has learned a structural causal model of the world. Instead, the benchmark generator supplies controlled interventions and matched base/intervened environments. This distinction is important: the evaluation borrows the practical idea of holding background factors fixed while changing an explicit variable, but its claims are about robustness and diagnostic evaluation under known interventions.

Distribution shift and evaluation documentation. Distribution-shift benchmarks such as WILDS emphasize that aggregate in-distribution performance can obscure robustness failures under changed conditions [Koh et al., 2021]. Ranking analyses of challenge benchmarks show that leaderboard conclusions can be sensitive to aggregation choices, test-case resampling, and the metric used to compare rankings [Maier-Hein et al., 2018, Perlitz et al., 2024]. Recent automated benchmark work similarly uses rank-correlation agreement between generated and human-labeled evaluations as evidence for evaluation reliability [Chen et al., 2024]. Dataset and model documentation work argues that evaluation artifacts should state their intended use, scope, limitations, and provenance [Gebru et al., 2021, Mitchell et al., 2019]. MapShift differs from MiniGrid-style procedural sampling by constructing matched base/intervened pairs for a fixed explored environment; from WILDS-style supervised shift benchmarks by evaluating embodied state acquired through reward-free exploration before intervention; and from causal RL by using known generator interventions for diagnostic robustness rather than estimating causal effects under confounding.

3 Controlled post-intervention protocol

Protocol. Let $\mathcal{E}$ denote a distribution over embodied environments, $\mathcal{F}$ the set of intervention families, and Σ_f the severity ladder for family f. A base environment $e \sim \mathcal{E}$ is sampled and exposed to the agent for reward-free exploration. During exploration, the system M receives observations and actions but no downstream task reward. We write the resulting internal state as

$$z_M = \mathrm{Explore}(M, e, T_{\mathrm{exp}}), \tag{1}$$

where z_M may be a recurrent hidden state, an explicit memory, a learned graph, or an empty state in the no-exploration ablation. After exploration, an intervention operator $I_f(\sigma)$ is applied, where $f \in \mathcal{F} = \{\mathrm{metric}, \mathrm{topology}, \mathrm{dynamics}, \mathrm{semantic}\}$ and $\sigma \in \Sigma_f$ is a discrete severity level. The resulting environment is

$$e' = I_f(\sigma)(e). \tag{2}$$

A task $q \sim Q_c(e, e')$ is then sampled from task class $c \in \mathcal{C}$, where $\mathcal{C}$ contains planning, inference, and adaptation. The evaluated system is scored using z_M from exploration in e while solving or inferring in e':

$$S(M; e, f, \sigma, c, q) = \mathrm{Perf}(M, z_M, e', q)\,. \tag{3}$$

The population target of the benchmark is therefore

$$\mathrm{CEP}(M) = \mathbb{E}_{e\sim\mathcal{E},\, f\sim\mathcal{F},\, \sigma\sim\Sigma_f,\, c\sim\mathcal{C},\, q\sim Q_c(e,e')}\left[S(M; e, f, \sigma, c, q)\right]. \tag{4}$$

In the paper, we do not treat this expectation as a single scalar leaderboard. The main estimand is family-wise:

$$\mathrm{CEP}_f(M) = \mathbb{E}_{e,\, \sigma\in\Sigma_f,\, c,\, q}\left[S(M; e, f, \sigma, c, q)\right], \tag{5}$$

because the benchmark is designed to separate metric, topology, dynamics, and semantic failure modes.

Operational intervention semantics. The generator can be viewed as a deterministic map $G(u, x)$, where u contains motif identity, sample seed, and layout randomness, and $x = (x_{\mathrm{metric}}, x_{\mathrm{topology}}, x_{\mathrm{dynamics}}, x_{\mathrm{semantic}})$ contains the configurable environment factors. A family intervention constructs

$$e' = G(u, x_{-f}, \phi_f(x_f, \sigma))\,, \tag{6}$$

where x_{-f} denotes the non-target factors and ϕ_f is the family-specific transformation. This is the sense in which the benchmark creates matched post-intervention pairs: the background motif and seed are held fixed, while one declared factor is changed and validators check that non-target factors remain invariant. We do not use the evaluation to infer causal effects from observational data, and we do not claim latent-unit counterfactual identification. The protocol is a controlled test-time shift with explicit intervention provenance.

Empirical estimator. The executable study estimates $\mathrm{CEP}_f(M)$ with a finite set of motifs, seeds, severities, task samples, and model seeds. Let $\mathcal{D}_f$ be the set of valid episode records for family f and baseline M, after task rejection. The reported family-wise score is

$$\widehat{\mathrm{CEP}}_f(M) = \frac{1}{|\mathcal{D}_f|} \sum_{r\in\mathcal{D}_f} s_r, \tag{7}$$

where s_r is the task-normalized primary score for episode record r.

Protocol comparison. Protocol sensitivity is measured by replacing pieces of the evaluation operator while holding the generated study grid fixed. Same-environment evaluation replaces e' with e at task execution time. No-exploration evaluation replaces z_M with an empty memory. Short- and long-horizon evaluations multiply task horizons by fixed factors. Rank changes are computed from the induced orderings

$$\pi_p(f) = \mathrm{rank}_M\left(\widehat{\mathrm{CEP}}_{f,p}(M)\right), \tag{8}$$

where p indexes the protocol variant. We summarize protocol sensitivity using Kendall τ, rank reversals, best-method changes, and family-wise rank changes.

Why post-intervention evaluation. The protocol separates two questions that are often conflated. Same-environment evaluation asks whether a system can exploit a representation in the environment it observed. Post-intervention evaluation asks whether that representation remains useful under a specified controlled change. This is close to the diagnostic question behind world-model evaluation [Ha and Schmidhuber, 2018, Hafner et al., 2020, Schrittwieser et al., 2020]: what structure did the model learn, and how does that structure fail?

4 Methodology

Executable study pipeline. The benchmark study is generated from versioned configuration files rather than manually curated episode lists. For each motif and sample seed, the generator writes a base environment, deterministic identifier, and manifest. Baselines run the configured reward-free exploration budget; learned methods train or load checkpoints keyed by baseline, environment identifier, model seed, and hyperparameter hash. Intervention operators and task samplers then produce the changed environments, task records, invariant checks, and rejection counts used by the health report.

Intervention construction. Table 1 summarizes the four intervention families. Each family changes one intended axis while preserving other axes when possible. Severity level 0 is the identity operator and is included as a sanity check; levels 1–3 increase the configured severity parameter. Severity is an ordinal within-family ladder, not a cross-family physical unit: severity 2 in topology is not claimed to be equally difficult to severity 2 in dynamics. The chosen ladders span small, medium, and large changes for each family and are interpreted through health checks for reference solvability, weak-baseline non-saturation, intervention isolation, and within-family severity response.

Table 1: Intervention methodology. Each family has a scalar severity parameter, preservation constraints, and expected failure modes used for benchmark health checks and interpretation.

Family	Severity parameter	Values	Preserved quantities
Metric	action-gain scale	1.0, 1.1, 1.25, 1.5	topology, semantic identities
Topology	blocker fraction	0.0, 0.08, 0.18, 0.32	semantic identities, action response
Dynamics	friction multiplier	1.0, 0.9, 0.75, 0.6	geometry, semantic identities
Semantic	remap fraction	0.0, 0.1, 0.25, 0.5	geometry, action response

Environment and observation model. The reference substrate uses 96×96 occupancy maps at 0.25m resolution. The action interface is semi-continuous navigation with a 0.25m forward primitive and 15° turn primitive. Observations are egocentric and local: the agent observes a 3.5m radius, 100° field of view, and semantic channels when visible. Exploration uses a single episode, uniform navigable starts, and no privileged global state for non-reference baselines.

Exploration budget. The exploration budget is held fixed at $T_{\text{exp}} = 800$ steps for all motifs and non-reference baselines. We do not normalize the budget by motif size or graph diameter because exploration insufficiency is itself part of the evaluated condition: a representation that works only after near-complete coverage should be distinguishable from one that remains useful under partial coverage. The benchmark health report therefore includes visited-cell ratios, visited-node ratios, path-length distributions, and weak-baseline saturation checks, so that under-exploration in larger motifs is visible rather than hidden by the protocol.

Figure 2 shows the matched-pair construction used by these interventions.

Intervention validation. After each intervention, validators check that the intended factor changed and declared invariants were preserved: metric shifts preserve topology and semantics, topology shifts preserve cue identities and action-response parameters, dynamics shifts preserve geometry, and semantic shifts preserve navigable geometry and reachability. These health checks are never model-facing; failures are reported before model results and block a valid release artifact.

Task sampling. The three task classes use shared start, goal, visibility, and horizon policies across comparable conditions. Planning covers shortest-path navigation, rerouting, changed dynamics, and

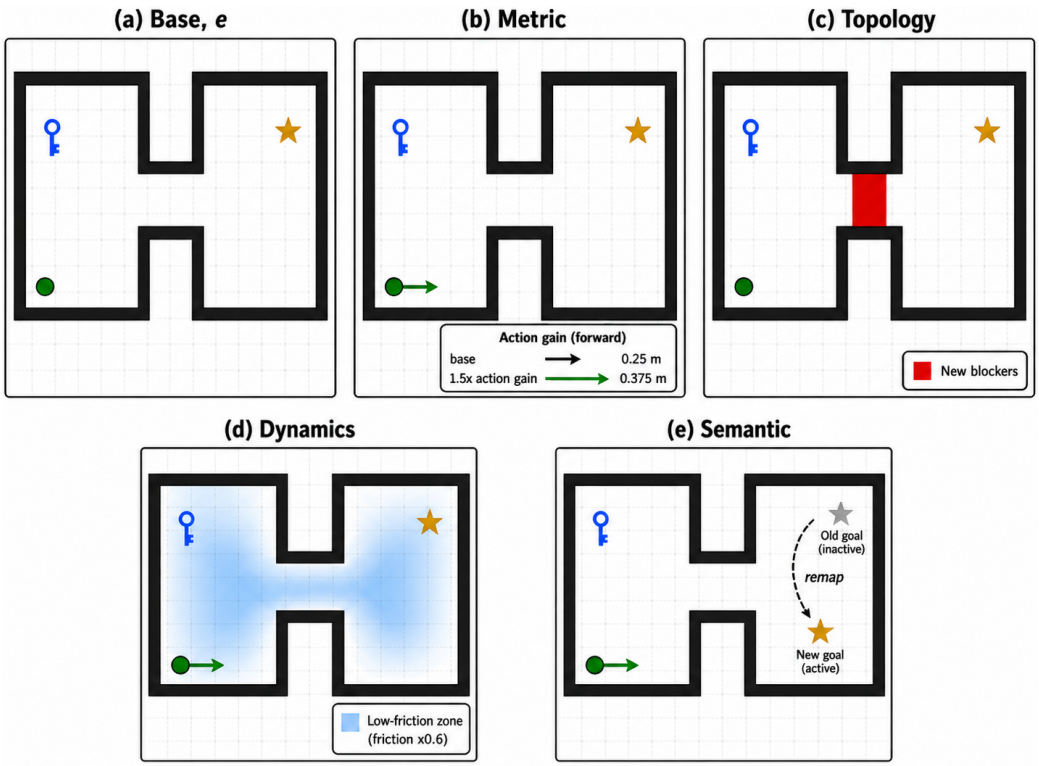

Figure 2: Matched intervention pairs in MapShift. A single base environment (a, two-room connector motif) is shown alongside the same environment after high-severity (level 3) interventions from each of the four families: metric (b, action-gain rescaling), topology (c, new corridor blockers), dynamics (d, low-friction zones), and semantic (e, goal-token remap). Background motif, seed, key location, and agent start are held fixed across panels; only the declared environment factor changes. Validators check that non-target factors remain invariant after each intervention. Severity-3 levels are shown for visibility; the full study uses severities 0, 1, 2, and 3 per family.

changed cues; inference covers topology-change detection, masked-region prediction, and counterfactual reachability; adaptation covers limited post-shift interaction followed by replanning. Planning horizons are 64, 96, and 128 steps; inference horizons are 16 and 32; adaptation budgets and horizons are 16, 32, and 64.

Task rejection and solvability. The sampler rejects tasks that cannot be solved by the privileged planning reference under the intended environment, tasks whose answer is unchanged in a way that makes the post-intervention query vacuous, and tasks whose shortest path or horizon makes the episode trivial. Rejections are not silently discarded: the artifact records rejection counts by reason and by benchmark cell. The health report therefore distinguishes a genuinely difficult benchmark cell from a cell that was under-covered because valid tasks could not be sampled.

Baseline evaluation. All non-reference baselines share $T_{\text{exp}} = 800$. The privileged post-intervention planning reference sees the changed environment as a solvability/planning diagnostic, not as a deployable method or universal upper bound. The same-environment reference is a protocol control: it asks what would be concluded if tasks were executed in the explored environment. The stale-map planner is a sensitivity floor that stores the pre-intervention map and never revises it after the intervention. Classical belief update stores occupancy/token bindings and updates from local post-shift observations; it receives no intervention label, privileged mismatch locations, or global post-intervention map. Learned baselines use fixed hyperparameters for pretrained graph, persistent memory, relational graph, and structured dynamics. Each record stores the baseline, run identifier, model seed, protocol, motif, split, family, severity, task class/type, and grouping key.

Weak heuristic floor. The weak heuristic is a meaningful but limited floor. It stores visited cells, visible token bindings, and shallow local traversability, then plans only over remembered local structure without global map repair. It intentionally fails under structural revision and stale semantic bindings, making it stronger than random guessing but too weak to saturate post-intervention reasoning. Full deterministic baseline behavior is specified in Appendix B.

Information access. Non-reference methods receive only their exploration trace, task specification, and internal memory; the global map and intervention label are reserved for explicit reference baselines.

Scoring. Planning episodes record success, observed path length, privileged-reference path length, normalized path efficiency, reference path gap, and post-intervention planning correctness. For a planning record r, path efficiency is

$$\eta_r = \begin{cases} \min(\ell_r^*/\ell_r, 1) & \text{if the agent succeeds and } \ell_r > 0, \\ 0 & \text{otherwise,} \end{cases} \tag{9}$$

where ℓ_r^* is the shortest feasible reference path length in the evaluation environment and ℓ_r is the observed executed length. Post-intervention planning accuracy is the success rate on planning tasks whose correct answer depends on the intervention, excluding identity shortest-path sanity tasks. These planning-reference quantities are kept separate from the mixed family primary score: family scores aggregate planning, inference, and adaptation components and are not normalized to make the privileged planning reference an upper bound.

Inference episodes record predicted and expected answers and are scored by task type. For masked-region prediction, the target is the changed family label or masked state category, and accuracy is exact-match:

$$A_{\text{mask}} = \frac{1}{N_{\text{mask}}} \sum_{r=1}^{N_{\text{mask}}} \mathbf{1}[\hat{y}_r = y_r]. \tag{10}$$

For post-intervention reachability queries, the target is the post-intervention node or room associated with a changed semantic token; the score is exact-match reachability/location accuracy,

$$A_{\text{reach}} = \frac{1}{N_{\text{reach}}} \sum_{r=1}^{N_{\text{reach}}} \mathbf{1}[\hat{g}_r = g'_r], \tag{11}$$

where g'_r is the correct target in e'. For topology-change detection, records have binary labels $y_r \in \{0, 1\}$ and model scores c_r. Probabilistic baselines may report calibrated confidence; deterministic baselines use $c_r = 1$ for a predicted change and $c_r = 0$ otherwise. AUROC is computed from average ranks with ties:

$$\text{AUROC} = \frac{\sum_{r:y_r=1} \text{rank}(c_r) - n_+(n_+ + 1)/2}{n_+ n_-}. \tag{12}$$

Adaptation episodes record a recovery curve over the allowed post-shift interaction budget and summarize it with adaptation sample efficiency. The family primary score is computed from task-relevant components and is used for ranking only after family-wise aggregation. Confidence intervals use 1000 bootstrap resamples at 95% confidence, grouped by environment/model-seed unit so that repeated task records from the same environment/model-seed combination are not treated as independent.

5 MapShift benchmark design

MapShift uses a partially observed navigation substrate with occupancy, transition, and semantic layers. The design exposes geometry, connectivity, dynamics, and semantics independently while remaining cheap enough for many seeds, protocol comparisons, and health checks. Its purpose is not to compete with richer simulators on realism or scale; it is to make intervention and validation machinery inspectable as a reference artifact for future wrappers. The release contains eight structural motifs: simple loop, two-room connector, branching chain, asymmetric multi-room chain, offset bottleneck, nested bottleneck, deceptive shortcut, and disconnected subregion. Splits are motif-level;

Table 2: MapShift release values. These values define the main study and artifact contract.

Component	Value	Component	Value
Map size	96×96	Exploration budget	$T_{\text{exp}} = 800$
Motifs	8	Splits	3 train, 2 val, 3 test
Families	metric/topology/dynamics/semantic	Severities	0, 1, 2, 3
Task classes	planning/inference/adaptation	Task samples	3 per class/cell
Trace-trained seeds	5 per learned baseline	Bootstrap	1000 grouped resamples

the three held-out test motifs support leakage-controlled protocol comparison but not broad motif-level generalization claims. The primary estimand is family- and protocol-level behavior over the configured study grid.

6 Benchmark health and validity checks

An evaluation benchmark should be evaluated before its models are. The MapShift artifact therefore generates a benchmark health report before writing model results, covering motif splits, family/severity/task coverage, task rejection, reference solvability, weak-baseline non-saturation, path and horizon distributions, severity magnitude, intervention isolation, and leakage.

Task rejection is counted after built-in resampling, so zero task rejections means the released grid contains valid, non-vacuous tasks in every motif/family/severity/task-class cell, not that no internal proposal was discarded. The severity-magnitude gate checks configured intervention magnitude, not model-score monotonicity: scalar intervention magnitude must be nondecreasing from severity 0 to 3, while empirical scores may be nonmonotone because the primary score mixes tasks and finite samples. Saturation would occur if the privileged reference and weak heuristic were both near-perfect, or if severity curves were flat because all methods solved or failed every cell; MapShift treats such cases as health warnings.

Leakage is reported in three categories. Fatal leakage corresponds to split or task leakage that invalidates a result and must be zero. Diagnostic warnings identify suspicious patterns requiring inspection. Benign reusable templates are explicitly labeled with rationale, such as shared query forms that do not reveal held-out environments. In the completed run, fatal leakage, leakage warnings, intervention-validator failures, severity-magnitude failures, and final-release task rejections are all zero, while reference solvability is 1.000.

7 Experimental program

Baselines. The study suite reports protocol/diagnostic references plus weak heuristic, classical belief update, pretrained graph, persistent memory, relational graph, structured dynamics, and stale-map planners. Same-environment compares conclusions under a different evaluation assumption and is not deployable. Learned methods optimize graph targets for edge existence, geometric cost, traversal cost, and semantic-token location; full hyperparameters and parameter counts are in Appendix B.

Model seeds and aggregation. The relational and structured-dynamics trace-trained baselines use five model seeds; the larger pretrained graph and persistent-memory rows are reported on 864 non-identity episodes. Hyperparameters are specified in versioned configs before final evaluation, validation motifs are reserved for pre-final choices, and test motifs are not used for hyperparameter selection. Summaries average over model seeds where applicable, and bootstrap resampling is grouped by environment/model-seed unit.

Study grid. The main study covers all eight motifs, all four intervention families, severities 0–3, planning/inference/adaptation tasks, and three task samples per class. It reports the primary CEP table as methods by intervention family, with severity-response diagnostics summarized in Appendix A and fully rendered in the artifact.

Protocol comparisons. To test whether evaluation assumptions change conclusions, we compare CEP to same-environment, no-exploration, short-horizon, and long-horizon variants. We report Kendall τ, rank reversals, best-method changes, family-wise rank changes, pooled-versus-familywise disagreement, and leave-one-family-out sensitivity.

8 Empirical calibration findings

The reference run calibrates the measurement infrastructure rather than establishing a definitive leaderboard. Health gates are clean: zero task rejections, zero fatal leakage, zero intervention-validator failures, zero severity-magnitude failures, reference solvability 1.000, and a non-saturating weak-heuristic health estimate of 0.721. Empirically, removing reward-free exploration changes all four family-wise method orderings (Kendall $\tau = 0.524$, five rank reversals), while same-environment versus CEP pooled rankings are stable for the full learned-probe roster. Aggregate stability therefore does not imply protocol irrelevance: the focused mechanism diagnostic changes the best method and separates stale reuse from explicit belief update.

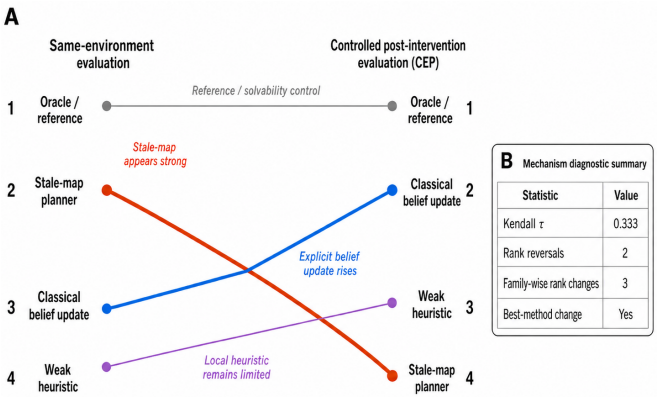

Statistic	Value
Kendall τ	0.333
Rank reversals	2
Family-wise rank changes	3
Best-method change	Yes

Figure 3: Mechanism diagnostic: CEP changes conclusions about stale reuse versus belief update. Panel A compares method ranks under same-environment evaluation and CEP; the figure's "oracle/reference" label denotes the privileged planning reference. Panel B reports Kendall $\tau = 0.333$, two rank reversals, three family-wise rank changes, and a best-method change.

In that diagnostic, same-environment evaluation rewards stale-map reuse because the remembered graph is still tested, while CEP invalidates it under topology blockers and semantic remaps. Belief update scores 0.689 and 0.838 on topology and semantic shifts, versus stale-map scores of 0.447 and 0.228; the privileged planning reference is a solvability/planning check, not an upper bound for every mixed score. Capacity controls move learned rows away from the floor without changing the main diagnosis: the 1.14M-parameter pretrained graph model improves over smaller trace-trained probes but remains below classical belief update, and a 1.4M-parameter persistent-memory run scores 0.313–0.354. These findings support MapShift as measurement infrastructure for post-intervention evaluation, not as an agent contribution.

9 Conclusion

MapShift frames post-intervention embodied planning as evaluation infrastructure, not an agent contribution. Its matched interventions, family-wise estimands, validators, and protocol comparisons test whether exploration-derived state remains useful after controlled changes. The calibration study shows that CEP separates stale reuse from explicit update and that same-environment, stale-memory, and no-exploration controls can mislead. More broadly, MapShift targets a recurring evaluation failure in embodied AI: aggregate task success can obscure whether an agent updated its world model, reused stale state, or exploited local heuristics.

References

Charles Beattie, Joel Z. Leibo, Denis Teplyashin, Tom Ward, Marcus Wainwright, Heinrich Küttler, Andrew Lefrancq, Simon Green, Victor Valdés, Amir Sadik, Julian Schrittwieser, Keith Anderson, Sarah York, Max Cant, Adam Cain, Adrian Bolton, Stephen Gaffney, Helen King, Demis Hassabis, Shane Legg, and Stig Petersen. DeepMind Lab. *arXiv preprint arXiv:1612.03801*, 2016.

Marc G. Bellemare, Yavar Naddaf, Joel Veness, and Michael Bowling. The arcade learning environment: An evaluation platform for general agents. *Journal of Artificial Intelligence Research*, 47: 253–279, 2013.

Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. OpenAI Gym. *arXiv preprint arXiv:1606.01540*, 2016.

Lars Buesing, Theophane Weber, Yori Zwols, Sébastien Racanière, Arthur Guez, Jean-Baptiste Lespiau, and Nicolas Heess. Woulda, coulda, shoulda: Counterfactually-guided policy search. In *International Conference on Learning Representations*, 2019.

Jianlyu Chen, Nan Wang, Chaofan Li, Bo Wang, Shitao Xiao, Han Xiao, Hao Liao, Defu Lian, and Zheng Liu. AIR-Bench: Automated heterogeneous information retrieval benchmark. *arXiv preprint arXiv:2412.13102*, 2024.

Maxime Chevalier-Boisvert, Bolun Dai, Mark Towers, Rodrigo de Lazcano, Lucas Willems, Salem Lahlou, Suman Pal, Pablo Samuel Castro, and Jordan Terry. Minigrid & miniworld: Modular & customizable reinforcement learning environments for goal-oriented tasks. *arXiv preprint arXiv:2306.13831*, 2023.

Karl Cobbe, Christopher Hesse, Jacob Hilton, and John Schulman. Leveraging procedural generation to benchmark reinforcement learning. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 2048–2056, 2020.

Matt Deitke, Eli VanderBilt, Alvaro Herrasti, Luca Weihs, Kiana Ehsani, Jordi Salvador, Winson Han, Eric Kolve, Ali Farhadi, Aniruddha Kembhavi, and Roozbeh Mottaghi. ProcTHOR: Large-scale embodied AI using procedural generation. In *Advances in Neural Information Processing Systems*, 2022.

Linxi Fan, Guanzhi Wang, Yunfan Jiang, Ajay Mandlekar, Yuncong Yang, Haoyi Zhu, Andrew Tang, De-An Huang, Yuke Zhu, and Anima Anandkumar. MineDojo: Building open-ended embodied agents with internet-scale knowledge. In *Advances in Neural Information Processing Systems*, 2022.

Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64 (12):86–92, 2021.

David Ha and Jürgen Schmidhuber. World models. In *Advances in Neural Information Processing Systems*, 2018.

Danijar Hafner. Benchmarking the spectrum of agent capabilities. In *International Conference on Learning Representations*, 2022.

Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. Dream to control: Learning behaviors by latent imagination. In *International Conference on Learning Representations*, 2020.

Chi Jin, Akshay Krishnamurthy, Max Simchowitz, and Tiancheng Yu. Reward-free exploration for reinforcement learning. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 4870–4879, 2020.

Leslie Pack Kaelbling, Michael L. Littman, and Anthony R. Cassandra. Planning and acting in partially observable stochastic domains. *Artificial Intelligence*, 101(1–2):99–134, 1998.

Pang Wei Koh, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, Akshay Balsubramani, Weihua Hu, Michihiro Yasunaga, Richard Lanas Phillips, Irena Gao, Tony Lee, Etienne David, Ian Stavness, Wei Guo, Ben Earnshaw, Imran S. Haque, Sara Beery, Jure Leskovec, Anshul Kundaje, Emma Pierson, Sergey Levine, Chelsea Finn, and Percy Liang. WILDS: A benchmark of in-the-wild distribution shifts. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 5637–5664, 2021.

Eric Kolve, Roozbeh Mottaghi, Winson Han, Eli VanderBilt, Luca Weihs, Alvaro Herrasti, Daniel Gordon, Yuke Zhu, Abhinav Gupta, and Ali Farhadi. AI2-THOR: An interactive 3d environment for visual AI. *arXiv preprint arXiv:1712.05474*, 2017.

Chaochao Lu, Bernhard Schölkopf, and José Miguel Hernández-Lobato. Deconfounding reinforcement learning in observational settings. *arXiv preprint arXiv:1812.10576*, 2018.

Lena Maier-Hein, Matthias Eisenmann, Annika Reinke, Sinan Onogur, Marko Stankovic, Pascal Scholz, Tal Arbel, Hrvoje Bogunović, Andrew P. Bradley, Aaron Carass, et al. Why rankings of biomedical image analysis competitions should be interpreted with care. *Nature Communications*, 9(1):5217, 2018.

Russell Mendonca, Oleh Rybkin, Kostas Daniilidis, Danijar Hafner, and Deepak Pathak. Discovering and achieving goals via world models. *arXiv preprint arXiv:2110.09514*, 2021.

Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model cards for model reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency*, pages 220–229, 2019.

Yotam Perlitz, Ariel Gera, Ofir Arviv, Asaf Yehudai, Elron Bandel, Eyal Shnarch, Michal Shmueli-Scheuer, and Leshem Choshen. Benchmark agreement testing done right: A guide for LLM benchmark evaluation. *arXiv preprint arXiv:2407.13696*, 2024.

Xavier Puig, Kevin Ra, Marko Boben, Jiaman Li, Tingwu Wang, Sanja Fidler, and Antonio Torralba. VirtualHome: Simulating household activities via programs. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 8494–8502, 2018.

Danilo Jimenez Rezende, Ivo Danihelka, George Papamakarios, Nan Rosemary Ke, Ray Jiang, Theophane Weber, Karol Gregor, Hamza Merzic, Fabio Viola, Jane Wang, Jovana Mitrovic, Frederic Besse, Ioannis Antonoglou, and Lars Buesing. Causally correct partial models for reinforcement learning. *arXiv preprint arXiv:2002.02836*, 2020.

Manolis Savva, Abhishek Kadian, Oleksandr Maksymets, Yili Zhao, Erik Wijmans, Bhavana Jain, Julian Straub, Jia Liu, Vladlen Koltun, Jitendra Malik, Devi Parikh, and Dhruv Batra. Habitat: A platform for embodied AI research. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 9339–9347, 2019.

Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap, and David Silver. Mastering atari, go, chess and shogi by planning with a learned model. *Nature*, 588: 604–609, 2020.

Ramanan Sekar, Oleh Rybkin, Kostas Daniilidis, Pieter Abbeel, Danijar Hafner, and Deepak Pathak. Planning to explore via self-supervised world models. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 8583–8592, 2020.

Mohit Shridhar, Jesse Thomason, Daniel Gordon, Yonatan Bisk, Winson Han, Roozbeh Mottaghi, Luke Zettlemoyer, and Dieter Fox. ALFRED: A benchmark for interpreting grounded instructions for everyday tasks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10740–10749, 2020.

Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2 edition, 2018.

Yuval Tassa, Yotam Doron, Alistair Muldal, Tom Erez, Yazhe Li, Diego de Las Casas, David Budden, Abbas Abdolmaleki, Josh Merel, Andrew Lefrancq, Timothy Lillicrap, and Martin Riedmiller. Deepmind control suite. *arXiv preprint arXiv:1801.00690*, 2018.

Archana Warrier, Dat Nguyen, Michelangelo Naim, Moksh Jain, Yichao Liang, Karen Schroeder, Cambridge Yang, Joshua B. Tenenbaum, Sebastian Vollmer, Kevin Ellis, and Zenna Tavares. Benchmarking world-model learning. *arXiv preprint arXiv:2510.19788*, 2025.

Appendix

A Additional result tables

Table A.1: Conclusion audit enabled by MapShift. Standard or underspecified protocols invite the left-hand conclusion; MapShift supports the middle diagnosis. The source column separates the full learned-probe run from the focused mechanism diagnostic.

Standard conclusion	Source	MapShift diagnosis	Why it matters
Same-env pooled order is enough.	Full run + mechanism diagnostic	Full-run pooled order is stable, but the belief-update diagnostic changes best method ($\tau = 0.333$).	Standard evaluation can reward stale state.
Stale maps are competitive.	Mechanism diagnostic	Same-env ranks stale-map too highly; CEP favors belief update under topology/semantic shifts.	Updating, not reuse, is the target competence.
Memory capacity is enough.	Capacity runs	Large persistent memory improves to 0.313–0.354 but remains below update-based references on most families.	Capacity does not substitute for explicit update.
No-exploration is close enough.	Full run	All four family orders change ($\tau = 0.524$, five reversals).	Exploration assumptions are load-bearing.

Table A.2: Benchmark health summary for the completed full run. Health checks are generated before model-result interpretation and included in the release artifact.

Check	Value
Environments and splits	8 motifs; 3 train, 2 validation, 3 test
Task coverage	1152 tasks; 0 task rejections
Reference solvability	1.000
Weak-baseline health estimate	0.721 success proxy; not primary CEP score
Fatal leakage / benign warnings	0 / 0
Intervention validator failures	0
Severity-magnitude failures	0

Table A.3: Protocol sensitivity in the full run. Ranking comparisons are computed from the same generated study grid while changing one evaluation assumption.

Comparison	Kendall τ	Rank reversals	Family changes
No exploration vs. reward-free exploration	0.524	5	4
Same-environment vs. CEP	1.000	0	2
Short vs. long horizon	1.000	0	3
Pooled vs. family-wise reporting	1.000	0	0

Table A.4: Protocol sensitivity in the deterministic mechanism diagnostic. This focused run compares the privileged planning reference, stale-map, weak-heuristic, and classical belief-update baselines.

Comparison	Kendall τ	Rank reversals	Family changes	Best changes
No exploration vs. reward-free	0.667	1	1	No
Same-env. vs. CEP	0.333	2	3	Yes
Pooled vs. family-wise	0.667	–	2	Yes

Table A.5: Severity-response summary from the full run. Entries are family primary scores for two calibration baselines across severities 0–3. The severity health gate validates monotone intervention magnitudes, not monotone empirical scores; full method-by-family severity data and rendered curves are included in the artifact.

Baseline	Family	S0	S1	S2	S3
Privileged planning ref.	Metric	0.521	0.486	0.471	0.443
Privileged planning ref.	Topology	0.521	0.552	0.750	0.740
Privileged planning ref.	Dynamics	0.521	0.760	0.792	0.781
Privileged planning ref.	Semantic	0.521	0.802	0.833	0.823
Weak heuristic	Metric	0.512	0.440	0.409	0.407
Weak heuristic	Topology	0.486	0.566	0.444	0.516
Weak heuristic	Dynamics	0.512	0.677	0.697	0.693
Weak heuristic	Semantic	0.486	0.122	0.122	0.141

Table A.6: Main family-wise CEP calibration results. Entries are mixed primary family scores, not planning-reference-normalized upper-bound scores. Rows are source-aligned: the belief-update row comes from the mechanism diagnostic; relational and structured-dynamics rows use the main five-seed configuration; and the pretrained graph and large persistent-memory rows are larger capacity runs with 864 non-identity evaluation episodes. The weak heuristic is a hand-coded floor with strong local spatial priors. The same-environment reference and privileged planning reference share deterministic shortest-path mechanics where applicable and are reported separately to support protocol-comparison interpretation. Full 95% grouped-bootstrap confidence intervals are included in the artifact JSON and rendered Markdown tables.

Method	Metric	Topology	Dynamics	Semantic
Privileged planning ref.	0.480	0.745	0.771	0.807
Same-env. reference	0.480	0.745	0.771	0.807
Weak heuristic	0.442	0.542	0.687	0.218
Classical belief update	0.513	0.689	0.699	0.838
Pretrained graph world model	0.295	0.327	0.400	0.401
Persistent memory	0.313	0.332	0.323	0.354
Relational graph	0.252	0.261	0.252	0.204
Structured dynamics	0.259	0.242	0.280	0.218

Table A.7: Family-wise mechanism diagnostic scores under CEP. Scores are mixed primary family scores aggregating planning, inference, and adaptation, so the privileged planning reference is not a universal upper bound. Planning-only reference path length and reference-gap metrics are tracked separately in the artifact. The diagnostic separates stale pre-intervention reuse, local heuristic priors, and explicit belief update.

Method	Metric	Topology	Dynamics	Semantic
Privileged planning ref.	0.480	0.745	0.771	0.807
Stale-map planner	0.438	0.447	0.771	0.228
Weak heuristic	0.442	0.542	0.687	0.218
Classical belief update	0.513	0.689	0.699	0.838

B Baseline metadata

Table A.8: Deterministic reference baselines. These methods have no trainable parameters and are included to calibrate solvability, protocol sensitivity, and weak-reference behavior.

Baseline	Access and stored state	Evaluation behavior	Training
Privileged post-intervention planning reference	Privileged access to the evaluation environment. Used as a solvability and planning-reference check, not as a deployable agent or mixed-score upper bound.	Deterministic shortest-path planning, exact inference answers, and reference adaptation curve when solvable.	0 params; none.
Same-environment reference condition	Protocol-control wrapper used only when execution is performed in the explored environment rather than the intervened environment.	Shares deterministic planning mechanics with the corresponding evaluation condition; interpreted as a protocol comparison, not a method upper bound.	0 params; none.
Weak heuristic	Stores visited cells, visited node ids, visible goal-token bindings, and limited local traversable structure. No global map revision.	Orthogonal visible-path planner over remembered structure; always predicts no topology change; answers semantic reachability with stale remembered token bindings; predicts masked family by task-type rule; capped heuristic adaptation curve.	0 params; none.
Stale-map planner	Stores the serialized pre-intervention environment from reward-free exploration and deliberately does not update it after intervention.	Plans shortest paths and answers semantic queries using the pre-intervention map, then is scored in the evaluation environment. Used as a diagnostic sensitivity floor for stale exploration-derived state.	0 params; none.
Classical belief update	Stores an occupancy map, edge structure, dynamics signature, and semantic-token bindings from exploration.	Uses local post-shift mismatch observations and the task/adaptation budget to update edges, costs, dynamics, and token bindings, then replans. It does not receive the intervention family label or global post-intervention map.	0 params; none.

Training source. Persistent-memory, relational-graph, and structured-dynamics learned baselines are trained separately for each base environment and model seed from the current exploration instance. The pretrained graph world model is a larger capacity-control run trained on generated graph instances before CEP evaluation. In all cases, evaluation-time tensors are built from the exploration-derived graph: node features contain normalized position, local free/blocked/semantic observation statistics, role indicators, and global dynamics features; pair labels supervise edge existence, normalized geometric path cost, and normalized traversal cost; node labels supervise semantic-token location.

Optimization and hyperparameters. Relational and structured-dynamics learned baselines use Adam with learning rate 0.01, validation fraction 0.2, and early-stopping patience 3. The large persistent-memory run uses 2048 memory slots, readout width 672, Adam learning rate 0.0005, 120 epochs, and patience 20. The pretrained graph world model uses Adam with learning rate 0.0005, 120 epochs, and patience 20. Hyperparameters are specified in configuration files before final evaluation; validation motifs are the only intended place for pre-final configuration choices. Test motifs are not used for hyperparameter selection. Checkpoints are keyed by baseline name, environment id, model seed, and hyperparameter hash, then reused across protocol variants.

Compute reporting. The smoke path is designed to run on CPU and completes in approximately 2–5 minutes for the smoke build itself; the one-command artifact audit takes longer because it also runs the test suite. The completed full run was executed on a single NVIDIA L4 GPU with 23GB

memory and CUDA enabled, using one Python process and cached learned-baseline checkpoints. It took about 11.5 wall-clock hours from the first family evaluation log entry to release-manifest completion. Preliminary/debug runs are not included in the main reported compute unless otherwise stated. The learned training jobs are small; wall time is dominated by task generation, planning/evaluation loops, bootstrap aggregation, and artifact rendering.

Table A.9: Learned baseline metadata. Parameter counts are representative realized trainable parameters for the release configuration; exact per-run counts, device metadata, checkpoints, and training summaries are written to run manifests.

Baseline	Stored representation and architecture	Supervision and optimization	Size
Pretrained graph world model	Larger graph message-passing model with hidden size 256, 6 message-passing steps, pair width 256, and dynamics width 128.	Pretrained on 4000 generated training graphs with 800 validation graphs, then evaluated under the same CEP task grid; 120 epochs; Adam lr 0.0005; patience 20; 864 evaluation episodes.	$\sim$1.14M params.
Persistent-memory world model	Builds the same graph data, but reads each node through learned memory slots. Slot attention stores persistent exploration-derived state.	Large run with 2048 memory slots and readout width 672; same self-supervised heads; 120 epochs; Adam lr 0.0005; validation fraction 0.2; patience 20; 864 evaluation episodes.	$\sim$1.4M params.
Relational graph world model	Uses node features and the explored adjacency matrix; applies a node encoder plus 2 message-passing steps over normalized adjacency.	Same self-supervised targets; 10 epochs; Adam lr 0.01; validation fraction 0.2; patience 3; 5 seeds.	396 params.
Structured-dynamics world model	Factorized geometry and dynamics encoders; edge and geometry heads use geometry factors, while traversal-cost prediction additionally receives global dynamics factors.	Same self-supervised targets; 10 epochs; Adam lr 0.01; validation fraction 0.2; patience 3; 5 seeds. Structured by restricting which latent factors feed each prediction head.	504 params.

C Limitations and future work

MapShift is not a deployment-realism benchmark and should not be interpreted as evidence of broad real-world robustness, language grounding, commonsense reasoning, or deployment readiness. The substrate simplifies perception, semantics, dynamics, and embodiment; severities are ordinal within family; and held-out motifs support leakage-controlled protocol comparison rather than broad motif-generalization claims. Future work should carry CEP and its health checks into richer 3D environments.

D Concrete porting sketch: MiniGrid intervention wrapper

This appendix makes the substrate-agnostic claim operational. A MiniGrid-style port would keep the CEP interface fixed while replacing the current map generator with a MiniGrid adapter. The adapter must implement five functions: `sample_base(seed, motif)`, `explore(env, budget)`, `apply_intervention(env, family, severity)`, `validate_pair(env, env')`, and `sample_tasks(env, env')`. The output schema should be identical to MapShift: split id, base environment id, intervention family, severity, task class, task id, final-release rejection reason if any, reference-solvability status, weak-baseline score, and per-episode records. The point is not to make MiniGrid look like MapShift, but to make a richer substrate satisfy the same measurement contract.

Table A.10: Concrete MiniGrid porting plan. The current MapShift validators mostly transfer, but visual/multimodal wrappers add substrate-specific checks.

Component	MiniGrid instantiation	Validators that transfer	Validators needing redesign
Base generator	Seeded rooms, doors, keys, walls, lava, and goal objects; motif id replaces map motif.	Split leakage, motif counts, reachable-state coverage, reference solvability.	Mission-template leakage if natural-language missions are used.
Metric shift	Change action cost, step budget, turn/forward cost, or grid-edge weights without changing walls or object identities.	Topology preservation, object identity preservation, path-cost monotonicity.	Cost semantics must be documented because MiniGrid has discrete actions rather than metric distance.
Topology shift	Add/remove doors, walls, blockers, or one-way passages while preserving object identities and start/goal distribution.	Connectivity delta, non-target semantic invariance, reference solvability, task rejection counts.	Door/key dependencies require checking both geometric and inventory-conditioned reachability.
Dynamics shift	Introduce slip, stochastic action failure, locked-door transition changes, or lava penalty changes.	Geometry and semantic preservation, severity-response checks, weak-baseline non-saturation.	Stochastic transitions require repeated rollouts or analytic transition matrices for reference scoring.
Semantic shift	Remap goal color/object type, key-door binding, or mission target while preserving geometry.	Geometry/topology preservation, target-binding change, non-vacuous task checks.	Language templates must be checked so the answer is not leaked by wording.
Tasks	Post-shift navigation, changed-door/key inference, masked-cell prediction, and limited interaction followed by replanning.	Family/task coverage, horizon distributions, final-release rejection accounting.	Observation wrappers must specify whether agents receive symbolic grids, pixels, missions, or both.

The same acceptance gates would apply before model results are interpreted: fatal leakage must be zero; final-release task rejection must be reported by cell; the privileged planning reference must solve the released tasks; the weak heuristic must not saturate; severity curves should be interpretable within family; and intervention validators must show that non-target factors are preserved. A ProcTHOR or Habitat port would use the same adapter contract, replacing MiniGrid's grid graph with a navigation graph or navmesh and replacing symbolic-object validators with object-instance, scene-layout, and rendering/visibility validators. This is why the current implementation is useful as a reference artifact: it provides executable tests for the wrapper, not merely a set of benchmark scores.

E Artifact contents

The release artifact contains versioned configs, split manifests, intervention and task recipes, benchmark health reports, raw per-episode records, family-wise tables, severity-response data, protocol-comparison outputs, rendered paper tables/figures, and top-level manifests. The artifact is an executable benchmark and generator, not a static dataset contribution. Generated environments and tasks can be regenerated from code and configs.

F Broader impact and ethics

MapShift is an evaluation artifact for embodied-planning research. It does not contain human-subject data, scraped web data, personally identifiable information, or deployed decision systems. The main risk is misuse of benchmark scores as evidence for real-world autonomy or robotics safety beyond the benchmark assumptions. The paper mitigates this by stating supported and unsupported claims, reporting benchmark health separately from model results, and emphasizing family-wise diagnostics over a single leaderboard score.

G NeurIPS Paper Checklist

1. **Claims.** Yes. The abstract and introduction state that MapShift is an executable post-intervention evaluation protocol and benchmark, with claims limited to controlled post-intervention embodied planning.
2. **Limitations.** Yes. Appendix C states that MapShift is not a deployment-realism benchmark and should not be interpreted as evidence of broad real-world robustness, language grounding, commonsense reasoning, or deployment readiness.
3. **Theory, assumptions, and proofs.** N/A. The paper formalizes an evaluation protocol and estimands but does not present theoretical theorems requiring proof.
4. **Experimental result reproducibility.** Yes. Methodology, artifact, and appendix sections describe versioned configs, task generation, baseline training, grouped bootstrap estimation, and release artifacts.
5. **Open access to data and code.** Yes. The submission includes an executable artifact containing the benchmark code, versioned configs, reproduction commands, generated outputs, raw per-episode records, health reports, and rendered result tables needed to reproduce the main study. The artifact includes both a smoke-test path and full-study reproduction instructions. The contribution is an executable artifact rather than a static hosted dataset.
6. **Experimental setting/details.** Yes. The paper specifies splits, exploration budget, intervention families and severities, task classes, baseline mechanisms, hyperparameters, model seeds, and aggregation.
7. **Experiment statistical significance.** Yes. The paper specifies 95% bootstrap confidence intervals with 1000 resamples grouped by environment/model-seed units.
8. **Experiments compute resources.** Yes. Appendix B reports the completed full-run setting: a single NVIDIA L4 GPU with CUDA enabled and approximately 11.5 wall-clock hours. It also reports the CPU smoke-test setting and clarifies that preliminary/debug runs are not included in the main reported compute.
9. **Code of ethics.** Yes. The work is designed as a controlled evaluation artifact and does not involve human-subject data, private data, or high-risk deployed systems.
10. **Broader impacts.** Yes. The appendix discusses the main risk: over-interpreting benchmark performance as real-world autonomy readiness.
11. **Safeguards.** N/A. The paper does not release a high-risk pretrained model or dual-use generative model.
12. **Licenses.** Yes. The released code is licensed under the MIT License; generated benchmark outputs and rendered result artifacts are licensed under CC BY 4.0; third-party dependencies and referenced tooling are listed in `THIRD_PARTY_LICENSES.md`.
13. **Assets.** Yes. The paper documents the released benchmark artifact, its generated outputs, limitations, and intended use. No static dataset is introduced.
14. **Crowdsourcing and human subjects.** N/A. The work uses no crowdsourcing and no human subjects.
15. **IRB approvals.** N/A. No human-subject research is conducted.
16. **Declaration of LLM usage.** N/A. LLMs are not a component of the benchmark methodology, baselines, or experimental results.